



Touchstone

*A stone you rub gold against to find out whether it is real.
An honest place to test whether a trading idea is.*

White paper · Version 1.2

August 2026

Published by **Zb International Co., Ltd.**

24/29 Moo 8, Sao Thong Hin Subdistrict,

Bang Yai District, Nonthaburi 11140, Thailand

Registration 0125567021983 · info@zb-international.com

Touchstone is educational simulation software. It places no orders, handles no money, and gives no investment advice. Simulated results are not real results. Please read the risk statement in section 13 before relying on anything in this document.

Contents

1	Summary
2	Why this exists — three accounts, and what each one cost
3	The problem: too many forces, one pair of eyes
4	What Touchstone is
5	How it is built
6	The reading layers — news, data, positioning, conditions
7	The Lab — testing an idea without paying for the lesson
8	The honesty machinery
9	The Journal — your own record, measured the same way
10	The paths you did not take
11	What Touchstone is not
12	Privacy, and what leaves your phone
13	Risk statement
14	Roadmap and how it will be decided

1 Summary

Touchstone is an educational simulation tool for people who trade gold and other volatile instruments in their own time. It does three things. It gathers what is moving the market right now — news, released economic data, large-fund positioning, and the state of the chart — and scores it in one place. It runs a set of trading bots on imaginary money, in public, and publishes their results whether they are good or bad. And it gives the reader a laboratory in which to test their own rules against a real year of prices, with the specific tests that reveal whether a good result was skill or luck.

It is built on an unusual premise for this industry: **that the most valuable thing a trading tool can tell you is that your idea does not work.** Most of the software in this space is designed to make a strategy look good. Touchstone is designed to find out. The difference shows up in what it measures — not just what a strategy returned, but whether it survives a different run of luck, a broker's costs, a change of market conditions, and, above all, the honest test of being judged on data that did not exist when it was tuned.

Underneath all of it sits one idea. A real account only ever shows you the path you took. Touchstone lets you run the ones you did not — backwards over a year that has happened, and forwards alongside the money you are actually trading, so that in three months the question "should I have done it differently?" has an answer instead of an opinion.

No real order is ever placed. No real money is ever handled. Nothing the user types leaves their device.

2 Why this exists — three accounts, and what each one cost

I trade in my free time. I am not an institution and I do not manage anyone's money. What follows is my own history, and I include it because the app came directly out of it.

This is not a success story. Every account below ended at zero. I am setting them down because each one failed for a different reason, and the reasons turned out to matter more than the returns.

The first: \$50 to \$1,500 in three months

Then the US election result arrived and it was gone. I had no framework for what a scheduled political event does to a leveraged position. I was not wrong about the market; I simply had not considered that a date on a calendar could end an account.

The second: \$450 to \$4,700 in six months

This one still makes me laugh, and it is the reason Touchstone exists at all. I was learning to backtest, using MetaTrader's EA editor. Some toggles were set in a way I did not understand, and the test ran against the live chart rather than history. The account was wiped out — not by a bad trade, not by the market, but by *trying to learn whether my idea was any good*.

Learning to test a strategy should not cost you your account. That sentence is the entire product.

The third: \$1,700 to \$12,000 in six months

Better strategies this time, built after sitting down and honestly cataloguing what had gone wrong before. And then a small piece of ego: one position was still red, and instead of taking the roughly \$17,000 that was available to me, I waited for it to turn green. Silver fell 42%. The account went with it.

That loss was not a market failure. The market did what markets do — a strong intentional move against crowded positioning, and profit-taking on top. The failure was that I wanted to be right more than I wanted to be paid, and I had no record in front of me arguing otherwise.

What the three have in common

Not one of them was a strategy problem. The first was an event I had not planned for. The second was a tool that punished me for trying to learn. The third was ego, unopposed by evidence.

Somewhere in that sequence I stopped believing that trading rewards cleverness, and started believing it rewards three duller things: **discipline, resilience, and endurance**. Discipline to follow the rule you set when you were calm. Resilience to look honestly at the wreckage and name what actually happened. Endurance to keep measuring long after the excitement has gone.

Touchstone is an attempt to build a tool for those three qualities rather than for the cleverness. One of the five bots in it is called Endurance, and that is not decoration.

3 The problem: too many forces, one pair of eyes

A part-time trader in gold, silver, an index or a crypto pair is standing in the path of forces that do not care about them and do not arrive politely one at a time:

- **Scheduled economic data** — inflation, employment, growth — where what matters is not the number but the gap between the number and what was expected.
- **Central bank policy and buying**, which for gold in particular has been a dominant and under-watched force.
- **Geopolitical tension and war**, which move safe havens on no schedule at all.
- **Political noise**, including a single social-media post from a head of state.
- **High-frequency and algorithmic flow**, operating at speeds no human competes with.
- **Liquidity grabs** — moves engineered to trigger the stops of people positioned exactly where you are positioned.
- **Macro cross-currents** — the dollar, real yields, oil — each pulling on the others.

Any one of these is followable. All of them, continuously, while holding a job, is not. The honest experience of the part-time trader is not "I lack information" — it is **"there is too much of it, arriving from too many places, and I cannot tell which of it mattered."**

And underneath that sits the harder problem. Even with perfect information, you still do not know whether your own method works, because you have never measured it in a way that could have told you it does not.

4 What Touchstone is

Touchstone is a mobile web application with four screens.

Desk	Five paper-trading bots, each with an imaginary \$10,000, updated every half hour. What each holds, what it has made or lost, and a plain-language description of what gold is currently doing — trend, volatility, position within its 20-day range, momentum, which session is active, and what is scheduled next.
News	The readings that feed the bots' score: headlines that moved the market and by how much, economic figures released and scored against forecast, and where large funds and gold ETFs are positioned.
Bot Lab	A full year of real half-hour prices, replayed on the reader's own device. Move a slider and the entire year recomputes. Save a configuration and it begins keeping a forward record.
Journal	The reader's own trades, entered by hand and measured with exactly the same statistics as the bots — including a direct comparison against a mechanical rule over the same days.

It is available in English, Thai, Hindi, Chinese and Vietnamese, works offline once opened, and installs to a phone's home screen without an app store.

5 How it is built

An automated job runs every thirty minutes. It fetches a real gold price, advances every bot by one step, scores the news, scores any economic data released that morning, checks large-fund positioning, and writes down what happened. Nothing is retroactive: what a bot knew at 9:30 is what it knew, and the record of it cannot be edited afterwards.

The trading engine, the strategy code and the full history live in a private repository. What is published is display output only — the numbers, the price archive, and a single page. This separation is deliberate and does two things: it keeps the exact rules and weights private — the reasoning is described in section 6, the arithmetic is not — and it means the public face cannot accidentally leak the machinery behind it.

Deterministic replay

The replay engine in the reader's browser reproduces the original engine's year-long run **to the cent** before any control in the Lab is trusted. This is verified after every change. It is what makes an unusual feature possible: because the same settings over the same candles always produce the same trades, a saved bot's forward record can be recomputed from scratch on the device, from the moment it was saved. Nothing needs to be stored, and nothing can be quietly revised. The record is derived, not remembered.

The archive

18,312 half-hour candles of spot gold, covering 29 July 2025 to the present day. Every backtest, every luck test and every forward record in the application is computed from this one archive, on the reader's own device. No calculation is sent to a server, because there is no server.

6 The reading layers — news, data, positioning, conditions

Four independent readings are taken every cycle and combined into a single score that the bots trade with. Each is deliberately bounded, so that no one layer can dominate the others.

News

Headlines from the last two hours are read and scored against vocabularies of gold-positive and gold-negative language, weighted by how recently each arrived. The strongest few are shown to the reader with their individual scores, most gold-positive first — so a reading built on one loud story is visibly different from one built on three that agree, and anything pulling the other way is visible rather than averaged away.

The vocabulary is a written list rather than a judgement: each phrase carries a fixed weight, positive or negative, and a headline's reading is the sum of whatever phrases it contains. Headlines naming the central bank, inflation or the jobs report are amplified, because those move gold hardest. Nothing is inferred and nothing is learned — the same headline always produces the same number, today and in five years.

The weakness of this is worth stating plainly. Matching phrases is not reading English. A headline saying "no rate cut coming" contains the phrase "rate cut", and is scored as though a cut had been announced. That is the honest limit

of the method. It is why news is the smallest of the three votes, and it is the reason the economic-data layer was built at all: a released number cannot be misread that way.

Economic data

Eighteen United States series — inflation, wages, employment, claims, retail sales, growth — are taken from the Federal Reserve's own database and compared against what forecasters expected. The *surprise* is what is scored, not the number: a rise in inflation that everyone predicted is not news, and a small one that nobody did can be.

Which way a surprise points is fixed by category rather than decided case by case. Hotter inflation and stronger growth weigh on gold, because both keep interest rates high and gold pays no interest. More slack — rising unemployment claims, a higher jobless rate — helps it, because that is what brings rates back down.

Each figure's push is scaled by how large the surprise is against thresholds set per series, then faded by age — full weight for four hours, half weight to twelve, nothing after that. A single release can move the score by at most 0.6; the dollar-and-yield backdrop by at most 0.2; the total is capped at 0.8. These caps exist so that one dramatic morning cannot own the score.

Every released figure is shown to the reader in plain language: what came out, what was expected, and which way that points for gold.

Positioning — where the large money is

Two public sources: the US futures regulator's weekly report of how large speculators are positioned, and the daily holdings of the world's largest gold ETF. Read together they answer a question a chart cannot: is the money that moves this market currently leaning in or leaning out? Real yields, the dollar index and oil are read alongside as a macro backdrop.

Fund positioning is read as a position *within its own multi-year range* rather than as an absolute number, because "a large number of contracts" means nothing without knowing what large is for this market. Near the top of that range the trade is crowded — and crowding is treated as a **warning rather than an endorsement**: everyone who wanted in is already in, so there is little money left to push it further. Near the bottom there is room to buy. Funds adding, or cutting, for several weeks running counts separately, as evidence that a move has real backing rather than one week's noise.

The ETF is read week on week, with one extra check worth naming: holdings falling while the price rises is a rally nobody is funding, and is marked down for it.

Two further rules exist to stop this layer being confidently wrong. A regulator's report older than a fortnight has its vote halved, because a stale snapshot is a guess wearing a number. And when oil spikes while the dollar is *also* rising — the shape a geopolitical shock usually takes — the layer declares a **crosswind** and stands down rather than choosing a side. "War is good for gold" is a slogan. A frightened world buys dollars too, and gold priced in dollars then has a bid and a headwind at the same moment; the 2022 invasion weeks are the plain example, gold jumping and then giving most of it back while the dollar climbed.

Conditions

Finally, a description of the market itself: whether price sits above or below its long-term average and by how much, whether volatility is high or calm against its own three-month median, where price sits within its 20-day range,

momentum over the week and the month, which trading session is open, and what is scheduled next.

How the readings become one number

News and economic data are added together, each bounded before the addition and the sum bounded after it. Words and numbers agreeing therefore push harder than either alone; words and numbers disagreeing partly cancel, which is the correct outcome — a market being told two things is a market with less of a lean, not more.

Positioning is deliberately kept out of that sum. The regulator's report is a Tuesday snapshot published on a Friday, so it is at least three days old by the time anyone reads it: old enough that it must never be allowed to open a trade. It does two quieter things instead. It adjusts how large a position is — roughly half to one and a half times the ordinary size — and it can refuse a trade outright when it sits strongly on the other side. A slow signal is permitted to size, and permitted to veto. It is not permitted to trigger.

The market conditions are not scored at all. They are shown, and they are recorded against every trade the bots take, so that the question "when does this work?" can be answered later from the record instead of from memory.

An honest note on how this scoring was built

These rules were designed and refined with the assistance of AI tools, and they are worth describing accurately. They are **not** a machine-learning model and they do not learn from their own results. They are explicit, deterministic, weighted rules — each source given a gravity according to its historical importance and its freshness, each bounded so it cannot dominate, and all of them written down. That is a weaker claim than "AI-powered", and a more useful one: a rule you can read is a rule you can argue with, and one whose failures you can find. Whether this scoring predicts anything at all is an open question, which is exactly why the application keeps a public report card on it.

What is published here, and what is not

The method is described above in full. The individual weights, the per-series thresholds, and the rules of the bots not offered in the Lab are not published. That is a boundary worth explaining rather than leaving to be guessed at. It is not secrecy about the reasoning: the application already shows every headline with its own reading, every released figure with its own push, and which positioning rule fired — several times an hour, to anyone who opens it. What is withheld is the arithmetic, not the argument. And the part of this project that would be hardest for anyone to copy is not arithmetic at all. It is a measured record, and a record cannot be started in the past.

7 The Lab — testing an idea without paying for the lesson

This is the part of Touchstone that exists because of the second account.

The Lab holds two strategy archetypes, each fully adjustable. **Scout** trades the rubber band: when price stretches away from its short-term average and starts snapping back, it takes the snap. **Blend** follows the flow: when a fast average crosses a slow one and price clears a recent high or low, it goes with the move and lets winners run.

Every control is a slider, and moving one recomputes the entire year in under a second on the device. Risk per trade, stop width, patience, profit target, whether to trade only with the longer trend — and, behind a fold, the numbers that

used to be welded into the strategy: how far back the average looks, how far price must stretch before acting, when the stop moves to breakeven, how tightly the trail follows.

They are called reference bots for a reason. **They are not the product; they are the control group.** Their purpose is to be a yardstick a reader's own thinking can be measured against, and a worked example of what testing a rule honestly looks like.

Nothing in the Lab can lose money. That is the whole point of the second account's story: the place where you find out whether an idea works should not be the place where you find out what it costs.

8 The honesty machinery

Any tool can show a backtest. What distinguishes Touchstone is the set of questions it asks after the backtest, each of which exists to catch a specific way that a good-looking result lies.

The luck test	Re-runs the entire year at every value of every slider. A <i>plateau</i> means the idea survives a range of settings and the exact number does not matter much. A <i>lonely spike</i> means a fluke was found. It reports how many settings were profitable and whether the current one sits one step away from a cliff.
A thousand alternative years	Re-draws the same kinds of trade at random, with replacement, a thousand times. This answers "could this year have lost money?" — which merely reordering the trades cannot, because reordering cannot change a total. It reports the typical worst drawdown, the unlucky one, and the share of alternative years that still made money.
Condition slices	Groups every trade by what the market was doing when it opened — trend, volatility, session. A strategy that only earns in one regime is leaning entirely on that regime, not an edge.
Over time	A rolling three-month result and four equal quarters. One yearly total can hide a strategy that died in March.
Broker realism	Spread, slippage and commission per ounce, plus overnight financing at the correct rate — charged only on nights actually held, with no weekend charge and the correct triple charge on the day the settlement date jumps the weekend.
The honesty split	The year cut in two. If the halves disagree, the setting is probably luck rather than skill.
Forward records	The one that matters most. From the day a bot is saved, it is replayed over every new candle since — the same test as a backtest, except it could not have been tuned to fit. Backtests flatter. Forward records tell the truth.

That last row deserves emphasis, because it is the answer to the deepest problem in this field. A backtest is a result computed on data that already existed when the settings were chosen. It is always, structurally, a description of a search rather than a discovery. A forward record cannot be. In this project's own testing, a configuration showing

+78% on a backtest produced a loss over the same period measured forward. That gap is the single most important number the application can show anyone, and it is shown by default.

9 The Journal — your own record, measured the same way

The bots are the control group. The reader's own trading is the experiment.

The Journal records real trades by hand — date, instrument, direction, size, entry, exit, costs, and, most importantly, **why**. It then measures them with exactly the same statistics used on the bots: win rate, profit factor, payoff ratio, worst losing run, worst drawdown, average per trade, and the honesty split.

It also replays a chosen bot over exactly the days the reader's log covers, from the same starting capital, and puts the two side by side. Not as a competition — the sizes and instruments differ, so it is a rough comparison rather than a fair race — but to answer one question that is otherwise unanswerable: **did my judgement add anything to a mechanical rule, or not?**

The note field is the part that matters most in the long run. It is the thing the reader will read in three months, when they have forgotten what they were thinking and can finally see it clearly. My third account would have survived if such a note had existed.

10 The paths you did not take

A real trading account shows you exactly one timeline: the one you chose. You place the trade, you set the stop where you set it, you close it when you close it — and whatever the alternative would have done is unknowable forever. You are left with a feeling about it, and feelings about counterfactuals are the least reliable evidence there is. Everyone remembers the trade they should have held. Almost nobody remembers the nine times holding would have been a disaster.

The future is decided by what we do in the present, and in a real account you only ever get to do one thing. Touchstone's purpose is to give you the others.

Backwards: the path you could have taken

A backtest is a path not taken, run in reverse. It asks what a rule would have done over a year that has already happened — including the parts of that year you were not watching, were not trading, or had already talked yourself out of.

Forwards: the path running alongside you

This is the more valuable half, and it is the reason saved bots exist. **You can trade your own money at your own broker, and have Touchstone run the alternative version of you at the same time, on the same market, day by day.**

Save the rule you almost followed. The tighter stop you decided against. The mechanical version of the discretion you are currently exercising. From that moment it keeps its own honest record, recomputed over every new candle — and because it was saved *before* the outcome, it cannot be tuned to flatter itself afterwards.

Three months later you are not choosing between two opinions. You are reading two records over the same days, in the same market, and one of them is the version of you that you did not become.

You cannot run your real account twice. This is the closest thing available: a second, third and fourth version of your decision, running in parallel, costing nothing, and keeping score.

And against your own hand

The Journal closes the loop. It replays a chosen rule over exactly the days your own logged trades cover, from the same starting capital, and puts the two side by side — your judgement against the mechanism, over identical market conditions. Several saved bots can be compared against each other in the same way, so a whole family of paths can be watched at once.

It is worth saying plainly what this is for. It is not to prove that a bot beats you, and quite often it will not. It is to convert *"I wonder whether I should have..."* from a private, permanent, unanswerable regret into a number you can look at.

The third account in section 2 is exactly this. I had roughly \$17,000 available and waited for one red position to turn green. Had a bot been running the rule I had written down when I was calm — take the money at the target, do not negotiate with a losing position — the comparison would have been sitting on my phone in black and white while I was busy arguing with myself. It would not have forced me to act. But I would have had to look at it.

11 What Touchstone is not

- **Not a signal service.** It never says buy or sell. The scores describe what has happened; they do not instruct.
- **Not investment advice.** We are not licensed advisers and we do not manage money.
- **Not a broker.** No orders are placed, no money is handled, no account can be funded.
- **Not a promise.** Most trading ideas that get measured honestly turn out not to work. Several in this project's own history did exactly that and were dropped.
- **Not a substitute for your own judgement.** The tool measures. Deciding remains yours.

12 Privacy, and what leaves your phone

There is no account and no sign-up. There are no cookies, no analytics, no advertising and no third-party trackers. Saved bots, the trade journal, the chosen language and every Lab setting are held in the browser's own storage on the reader's device. We never receive them, cannot read them and cannot recover them.

The application reads exactly one thing from the device: the clock and time zone, so that dates, sessions and "updated twelve minutes ago" are correct where the reader is. This needs no permission and is never transmitted. It cannot access location, contacts, camera, microphone or files.

Two requests leave the device: the scoreboard data and the price archive, both from our own site. There are no third-party scripts and no downloaded fonts — the application uses the typefaces already on the device, which is also why

it works with no signal.

13 Risk statement

Trading is risky and most people who attempt it lose money. Trading gold, silver, indices, cryptocurrencies, contracts for difference and other leveraged products carries a high risk of loss. With some brokers you can lose more than you deposit. You should not trade with money you cannot afford to lose entirely.

Simulated results are not real results. Every figure produced by this application is a simulation. It charges no spread or financing unless the reader sets one; it cannot suffer a rejected order, a requote, a gap through a stop, or a broker outage; and it is computed with hindsight over data that already exists. Real execution differs, and it differs most in exactly the fast conditions where it matters most.

Past performance never predicts future results. This is not a formality. A strategy that worked for a year can stop working permanently, and it will not announce it.

Nothing here is advice. Nothing in this application or this document is investment, financial, legal or tax advice, or a recommendation to buy or sell anything. Zb International Co., Ltd. is not a licensed investment adviser, fund manager or broker. Any decision taken in a real market is entirely the reader's own.

The data may be wrong. Prices, news and economic figures come from third-party sources and may be delayed, incomplete or incorrect. The application is provided "as is" with no warranty of availability, accuracy or continuity.

The author's own record is a warning, not a credential. The accounts described in section 2 all ended at zero. They are included to explain why the tool exists, not to suggest that anyone should expect similar gains — or to imply that the losses were unusual. They were not.

If you are new: the responsible use of this tool is to find out that an idea does not work, before that discovery is made with money.

14 Roadmap and how it will be decided

A measuring era began on 15 August 2026. From that date the bots run under their current rules, with the economic-data layer live, and every data morning is recorded: the score when the bots first knew, and where gold stood four and twenty-four hours later.

Six to twelve months of that record decides what happens next. If it shows that these readings predict something, further work is justified — more instruments, shorter timeframes, alerts, deeper tests. If it shows they do not, that will be published here too, and the tool will carry on being a good place to discover that a strategy does not work, which is worth knowing before it costs money.

What will not change: the honest measurement is the point of the project, and a tool that hid whether a strategy was luck would be the opposite of what this is for.

Contact

Zb International Co., Ltd., 24/29 Moo 8, Sao Thong Hin Subdistrict, Bang Yai District, Nonthaburi 11140, Thailand. Registration 0125567021983. info@zb-international.com

© 2026 Zb International Co., Ltd. All rights reserved. Touchstone, its bot rules, texts and designs are the property of Zb International Co., Ltd. This document is offered in English; where a translation and the English version disagree, the English version applies.

